
Fuzz Testing

Aaron Yang

Questions

Question 1 (10 pts): What is Fuzz Testing (basic idea) and its challenges? Which method does AFL use to solve the problems? (Hint: see its documentation)

- Fuzz testing is a software testing technique that involves sending a large number of random and invalid inputs or data to a program to detect vulnerabilities, bugs, and crashes. The basic idea behind fuzz testing is to identify unexpected and unhandled conditions, which can be exploited by attackers to compromise the security of the software. The main challenges of fuzz testing include generating inputs that are diverse, meaningful, and effective at triggering faults, as well as handling the large number of inputs and the crashes they may cause. AFL instruments the target program to measure code coverage, applying mutation operators to the initial input corpus, and using a feedback-driven genetic algorithm to guide the generation of new inputs.

Question 2 (10 pts): Briefly explain the overall algorithm of AFL approach. (Hint: see its documentation)

- The AFL approach involves instrumenting the target program to measure code coverage, using an initial set of inputs to generate a corpus, executing the program, selecting the most promising inputs based on their fitness, and repeating the mutation and selection process. AFL uses a feedback-driven genetic algorithm to guide input generation towards achieving higher code coverage.

Questions

Question 3 (10 pts): How does AFL measure the code coverage during the fuzzing process? List formula it uses for calculation. (Hint: see its technical details file)

- AFL measures the code coverage during the fuzzing process by instrumenting the target program with a code coverage feedback mechanism. AFL counts the number of basic blocks covered by each input and calculates the edge coverage and path coverage using the following formulas:
- Edge coverage: $\text{number of edges covered} / \text{total number of edges}$
- Path coverage: $\text{number of unique paths covered} / \text{total number of unique paths}$
- AFL also uses a feedback-driven genetic algorithm to guide the generation of new inputs towards achieving higher code coverage.

Question 4 (10 pts): What is fork server? Why does AFL use it during the fuzzing?

- A fork server is a mechanism used by AFL to improve the efficiency and stability of the fuzzing process. A fork server creates a child process of the target program that can be reused for multiple fuzzing iterations, instead of starting a new instance of the program for each input. This approach avoids the overhead of repeatedly loading and initializing the program, and it also isolates the target program from crashes and other errors caused by the inputs. The fork server can also detect crashes and other anomalies in the child process and report them to the AFL main process for analysis.

Extra Credit Question

One whitebox fuzzing approach that could be used to solve the problem is symbolic execution. On symbolic execution, the program is executed symbolically with input variables represented as symbols. The program path constraints are collected during the symbolic execution, which are then solved to generate inputs that can satisfy the constraints and explore new paths.

To apply symbolic execution to the given code example, the following steps could be taken:



1. Represent the input variable x symbolically.
2. Execute the program symbolically with x as input and collect path constraints during execution.
3. Use a constraint solver to generate inputs that satisfy the path constraints.
4. Execute the program with the generated input to explore new paths and repeat the above steps.

By using symbolic execution, the whitebox approach can generate a particular input that triggers the error.

Screenshots

It is zero or not - 7:32

The screen - 7:35



Screenshots

Indication of the crash file

The screenshot displays a desktop environment with a background of mathematical formulas. On the left, there is a dock with various application icons including Recycle Bin, Zoom, Google Chrome, Oracle VM VirtualBox, Riot Client, VMware Workstation 17 Player, Steam, Microsoft Edge, Discord, Atbm, and League of Legends. The main window is a terminal titled "Ubuntu 64-bit - VMware Workstation 17 Player (Non-commercial use only)". The terminal shows the output of the AFL-fuzz command, indicating a crash file has been found. The output is as follows:

```
aaron@aaron-virtual-machine: ~/Downloads/AFL-master
Exec timing : 1459 to 1459 us (average: 1717 us)
[*] No -t option specified, so I'll use an exec timeout of 20 ms.
[+] All set and ready to roll!

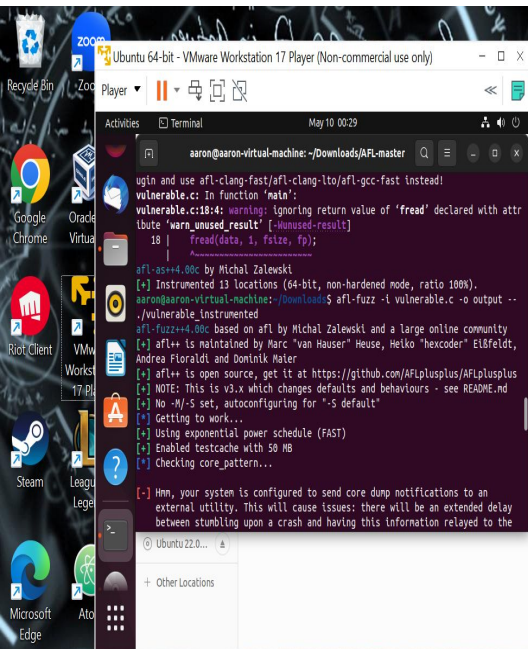
american fuzzy lop ++4.00c {default} (...loads/AFL-master/test-instr) [fast]
process timing
  run time : 0 days, 0 hrs, 9 min, 10 sec
  last new find : none yet (odd, check syntax!)
  last saved crash : none seen yet
  last saved hang : none seen yet
overall results
  cycles done : 3
  corpus count : 44
  saved crashes : 0
  saved hangs : 0
cycle coverage
  now processing : 1.166 (2.3%)
  runs timed out : 0 (0.00%)
map coverage
  map density : 0.00% / 0.00%
  count coverage : 1.00 bits/tuple
stage progress
  now trying : splice 15
  stage execs : 68/146 (46.58%)
  total execs : 556k
  exec speed : 830.3/sec
findings in depth
  favored items : 1 (2.27%)
  new edges on : 1 (2.27%)
  total crashes : 0 (0 saved)
  total tmouts : 0 (0 saved)
fuzzing strategy yields
  bit flips : disabled (default, enable with -D)
  byte flips : disabled (default, enable with -D)
  arithmetics : disabled (default, enable with -D)
  known ints : disabled (default, enable with -D)
  dictionary : n/a
  havoc/splice : 0/194k, 0/361k
  py/custom/rq : unused, unused, unused
  trln/eff : 33.33%/1, disabled
item geometry
  levels : 1
  pending : 0
  pend fav : 0
  own finds : 0
  imported : 0
  stability : 100.00%
[cpu000: 50%]

+++ Testing aborted by user +++
[+] We're done here. Have a nice day!


aaron@aaron-virtual-machine:~/Downloads/AFL-master$ afl-fuzz -i ./testcases -o ./output /home/aaron/Downloads/AFL^C
aaron@aaron-virtual-machine:~/Downloads/AFL-master$ ls ./output/default/crashes
aaron@aaron-virtual-machine:~/Downloads/AFL-master$ ls ./output/default
cmdline_crashes fuzz_bitmap fuzzer_setup fuzzer_stats hangs plot_data queue
aaron@aaron-virtual-machine:~/Downloads/AFL-master$
```

The terminal window also shows a status bar at the bottom with the text "cpu000: 50%". The desktop environment includes a taskbar at the bottom with a search bar and various application icons. The system clock in the bottom right corner shows 8:03 PM on 5/6/2023.


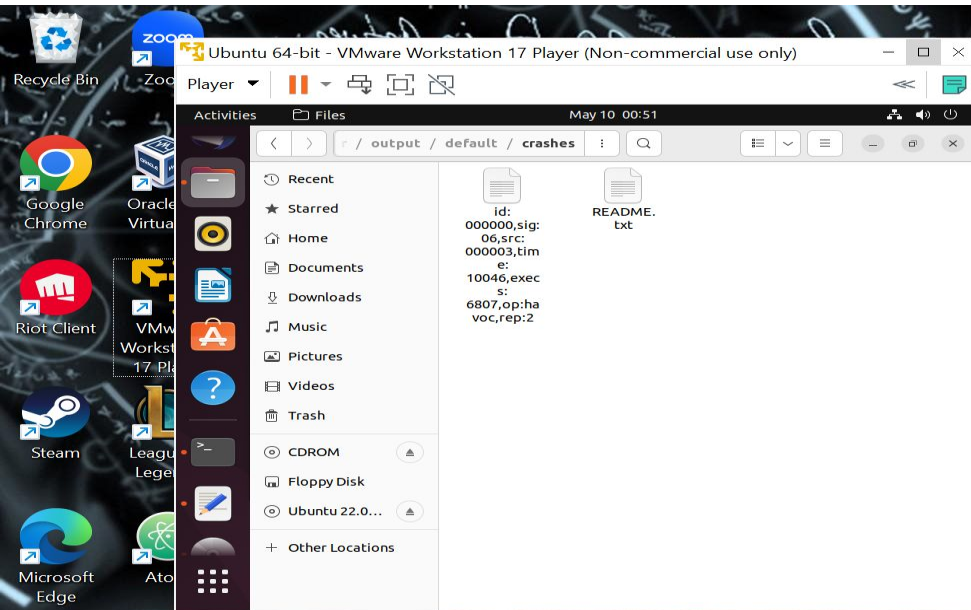
More Screenshots of Last Step



Input File



Crash File



Summary/Lab Report

First step indicates that we have to compile the test cases using this command `gcc -o ./test-instr ./test-instr.c` then `./test-instr` to see the inputs and outputs. If you get put a value that is not 0, it will return this is not zero but if zero, then say it is zero.

Second step indicates that we have to set up the fuzzing process. Using gcc to compile `afl-gcc -o ./test-instr ./test-instr.c`, you get to instrument and the instrumentation of the process. I created an output folder named output to store my crashes in there. After this I ran into a problem until I found my path executable to directly compile and fuzz it. I used this to fuzz it `afl-fuzz -i ./testcases -o ./output /home/aaron/Downloads/AFL-master/test-instr @@` with my path in there. We tested the cases and let it run for a few cycles.

Third step indicates that we have to find the crash file so it seems that it is stored in output so we did `ls ./output`. We examine the crashing test case and its content. It was stored in the default area so we access the crashes by `ls ./output/default/crashes` in order to find the crash file. We run with the test-instr program. `./test-instr < ./output/default/crashes/id:000000,op:havoc,rep:32`. It would cause a segmentation fault.

Last Step

For the last step, we want to compile the program of `vulnerable.c` using `afl-gcc -o vulnerable_instrumented vulnerable.c`.

Prepare the initial test input:

- Create a directory to store the initial test case files (e.g., "input").
- Place your initial test input file (e.g., "input_fuzz.txt") in the input directory.
- 2. Set up the AFL fuzzing command:
 - Open a terminal and navigate to the directory containing the compiled program and the input directory.
 - Run the AFL fuzzing engine with the following command: `afl-fuzz -i input -o output -- ./vulnerable_instrumented @@`
 - i. Replace "input" with the actual path to the input directory.
 - ii. Replace "output" with the desired path for AFL to store its findings.
 - iii. "vulnerable_instrumented" is the name of the compiled program.
 - Monitor the AFL execution status:
 - i. AFL will mutate the initial test input files and feed them to the program.
 - ii. It will monitor code coverage, detect crashes, and explore different program paths.
 - iii. The AFL execution status will be displayed in the terminal, providing information about the progress, number of iterations, unique crashes found, and other relevant details.
 - Analyze the output:
 - i. If AFL discovers crashes or abnormal behaviors during the fuzzing process, it will save the input that triggered the issue in the output directory.
 - ii. Analyze the crash inputs in the output directory to understand the vulnerabilities in the program.